

DI- FCT/UNL

May 8th, 2025**Database Systems****Test-1 2024/25****Duration: 2 hours (limited information)****Group 1 (each question is worth 2,5 values out of 20)**

Consider part of a database for a mobile phone operator that provides roaming services

Persons({personID,name,fiscalNR,countryID,...}) Countries({countryID,countryName,code})

Devices({deviceID,number,operatorID,personID}) Operators(operatorID,operatorName,countryID)

Registration(regID,deviceID,operatorID,start,end) Calls(callID,regCaller,regDest,time,duration,cost)

Persons (personID as primary key) stores user data and includes a foreign key countryID referencing **Countries** (countryID as primary key), which holds country names and dial codes. **Operators** (operatorID as primary key) also reference **Countries** to indicate their home country. **Devices** (deviceID as primary key) are foreign keys to **Persons** (personID) and **Operators** (operatorID), indicating ownership and issuing operator of the device contract with the operator. **Registration** (regID as primary key) logs device roaming sessions using foreign keys deviceID and operatorID, where start and end are timestamps (end might be null if the roaming session is active). **Calls** (callID as primary key) track roaming calls, with foreign keys regCaller and regDest to caller and destination of a call, referencing **Registration**, and include details like time, duration (integer in seconds), and cost. For each of these tables there is a secondary B+ tree (non-clustering) index on the primary key attribute(s), created with the column order indicated in the tables. Additionally, there is a composite secondary index on Calls(regCaller,regDest).

The adopted DBMS uses blocks of 4KiB (4096 bytes). The records of all tables have a variable size, and at any given time, the Persons table has 100,000 tuples (tuple size 108 bytes, total 2640 disk blocks), and Devices 200,000 tuples (62 bytes/3028 blocks), the Countries table 195 tuples (58 bytes/3 disk blocks), the Operators table 740 tuples (58 bytes, total 11 disk blocks), the Registration table has 857,000 tuples (28 bytes/5852 blocks), and the Calls table 10,000,000 tuples (28 bytes per tuple/68,360 disk blocks).

A B+ tree node can contain about 100 search keys, and it is known that a seek time is 10ms and the transfer time of a block t_T is 1ms, while the memory only holds 100 blocks.

Note: In this group, whenever examples are requested, these must be exclusively about this database, Additionally, all the answers must contain a brief justification.

1 a) Present two execution plans for the following SQL query, briefly justifying which of the plans has the least cost in the given database.

```
SELECT personID, name, deviceID, number
FROM Persons NATURAL INNER JOIN Devices NATURAL INNER JOIN Countries
WHERE countryName = 'Portugal' AND name BETWEEN 'Carla' AND 'Carlos'
```

1 b) Present a query for which the use of hash static file organization would be helpful.

1 c) Indicate whether in your opinion the SQL query below can benefit from the existing indices.
Would you introduce extra indices for optimizing the query?


```
SELECT * FROM Persons NATURAL INNER JOIN Devices
NATURAL INNER JOIN Registrations INNER JOIN Calls ON (regID = regCaller)
WHERE personID = 74578
```

1 d) Explain if a bitmap index could improve the linear scan execution of the following query:

```
SELECT SUM(cost) FROM Calls
WHERE duration <= 1000
```

(empty page – any content will not be graded)

1 e) Consider the execution plan presented in the figure. Determine the least cost of this plan knowing that there are 4 operators in Portugal, and that in Portugal there are 10.000 registered devices. Assume that the height of B+-tree for the non-clustering index of Countries is 2. Devices is used as probe relation.



Group 2 (each question is worth 2,5 values out of 20)

Note: The response to each of the items in this group cannot, under any circumstances, exceed one page.

- 2 a)** Explain why RAID LEVEL 6 is becoming more and more relevant, especially for disks with large capacities.

2 b) Explain why it is needed a NULL bitmap and a DELETED bitmap in bitmap indices.

2 c) If recursive partitioning is not required, the hash-join algorithm has a worst-case complexity of $3(b_r+b_s) + 4*n_h$ block transfers + $2(\lceil b_r/b_b \rceil + \lceil b_s/b_b \rceil) + 2*n_h$ seeks, when b_b disk blocks are used as blocking unit in the partitioning phase and n_h is the total number of partitions. Explain how the formula is derived (you may ignore the $4*n_h$ block transfers cost).

PSEUDO-CODE FOR HASH-JOIN

1. Partition the relation s using hashing function h . When partitioning a relation, one block of memory is reserved as the output buffer for each partition.
 2. Partition r similarly.
 3. For each i where $0 \leq i < n_h$ (where n_h is the total number of partitions):
 - (a) Load partition s_i into memory and build an in-memory hash index on it using the join attribute. This hash index uses a different hash function than the earlier one h .
 - (b) Read the tuples in partition r_i from the disk one by one. For each tuple t_r locate each matching tuple t_s in s_i using the in-memory hash index. Output the concatenation of their attributes.

Relation s is called the **build** input and r is called the **probe** input.

THE END

(SCRAP PAPER)